

# **Generative Models & Determining Trained**

Presenter: Kim Seung Hwan (overnap@khu.ac.kr)



“An astronaut riding a horse in a photorealistic style”



“An astronaut riding a horse as a pencil drawing”

OpenAI DALL-E 2

# Goal

- Determining if a model has trained on some data
- This is important in two ways as DNNs are becoming commercialized
  1. Copyright aspect; even more important in generative models
  2. Metric aspect; the model's training efficiency, generalization performance, and the value of new data for it
- It seems worth, but is it possible?

# Goal

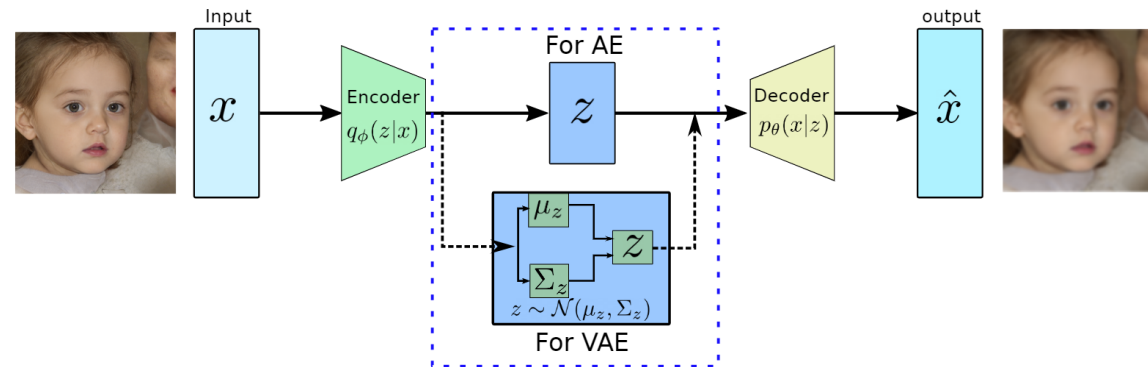
- My thoughts:
- In supervised learning, a model is forced to represent training data
- The model is optimized for the training data, not the true distribution
- The ideal generalization is difficult to achieve empirically
- It will be able to determine under appropriate constraints (but so what?)

# Progress

- Studying generative models, especially VAE and Diffusion Model
- Naïve experiments and analysis attempts on simplified problem

# VAE

- Variational Auto Encoder
- Let the latent of AE be indeterministic
- This will be some distribution (e.g. Gaussian RV)
- For encoder  $q(z|x)$  and decoder  $p(x|z)$ , we want to know  $p(z|x)$
- Indirectly perform MLE by applying variational inference ideas



# VAE

$$\log p(\mathbf{x}) = \log p(\mathbf{x}) \int q_{\phi}(\mathbf{z}|\mathbf{x}) d\mathbf{z} \quad \text{(Multiply by } 1 = \int q_{\phi}(\mathbf{z}|\mathbf{x}) d\mathbf{z} \text{)} \quad (9)$$

$$= \int q_{\phi}(\mathbf{z}|\mathbf{x}) (\log p(\mathbf{x})) d\mathbf{z} \quad \text{(Bring evidence into integral)} \quad (10)$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x})] \quad \text{(Definition of Expectation)} \quad (11)$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \quad \text{(Apply Equation 2)} \quad (12)$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z}) q_{\phi}(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x}) q_{\phi}(\mathbf{z}|\mathbf{x})} \right] \quad \text{(Multiply by } 1 = \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \text{)} \quad (13)$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right] + \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \quad \text{(Split the Expectation)} \quad (14)$$

$$= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right] + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}|\mathbf{x})) \quad \text{(Definition of KL Divergence)} \quad (15)$$

$$\geq \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right] \quad \leftarrow \text{ELBO} \quad \text{(KL Divergence always } \geq 0 \text{)} \quad (16)$$

# VAE

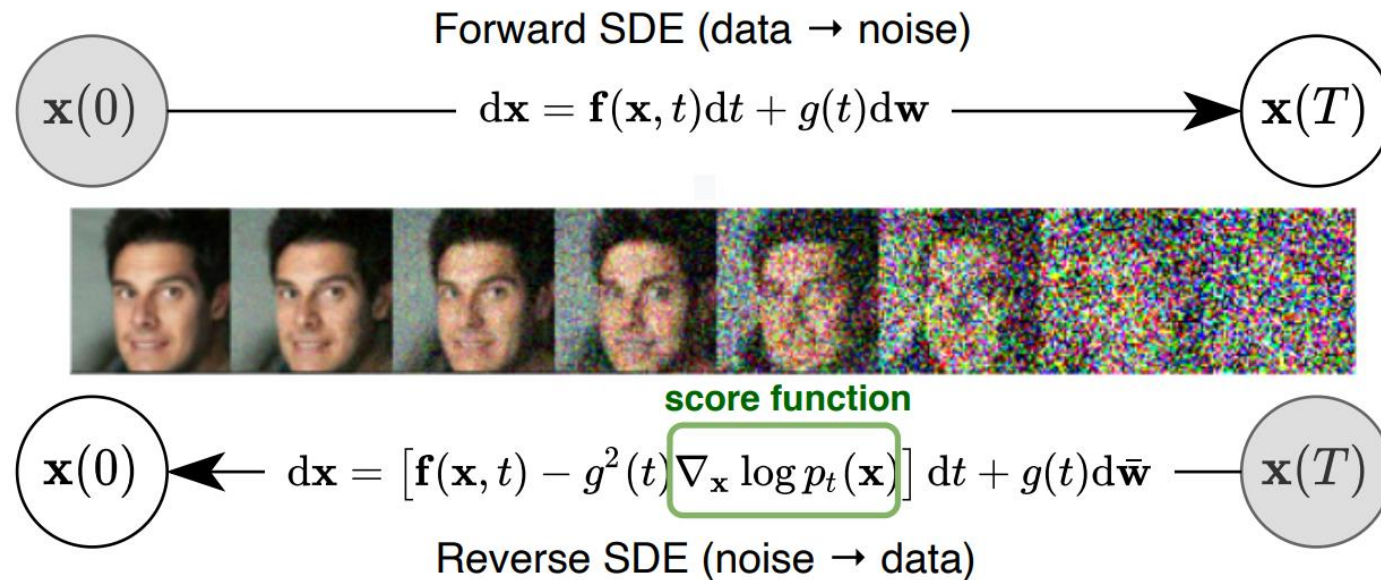
- “Increasing ELBO achieves MLE and posterior matching simultaneously”?
- What does MLE mean? Is it right to only grow the likelihood on training data?
- May fail to reduce true posterior matching term:
- When increasing ELBO, decreasing p.m. term is only guaranteed when  $\theta$  is fixed!

$$\begin{aligned} \theta \log p(\mathbf{x}) &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right] + \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}|\mathbf{x}))}_{\text{true posterior matching term}} \quad \phi \\ &\geq \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right]}_{\text{ELBO}} = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}))}_{\text{prior matching term}} \end{aligned}$$



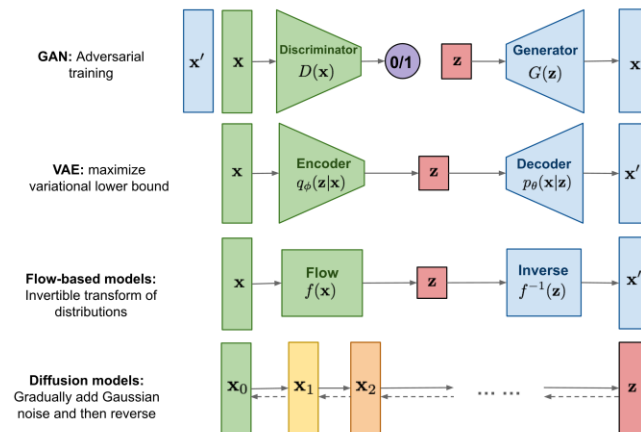
# Diffusion Model

- Adding small Gaussians to a sample to make it a full Gaussian
- It requires a lot of steps






# Diffusion Model

- Generative models convert data  $x$  to latent data  $z$  and then reconstruct them back to  $x$  or something similar
- Diffusion model is one of the score-based generative model
- The score is the gradient-log of the true data distribution  $\nabla_x \log p(x)$
- There is a lot of math involved, but it is hard to know the meaning...
- Let us approach it from a different perspective and then come back




# Diffusion Model

- A hierarchical structure in which several VAEs are connected
- This is called HVAE and is a generalization of VAE
- Like DNNs, HVAEs are likely to have more representation ability



# Diffusion Model

- But it is too free to model; Let us set each VAE to be a Markov chain
- Then, given the original data  $x = z_0$ , for each time  $t$  and latent  $z_t$ , it can be modeled with encoders  $q(z_t|z_{t-1})$  and decoders  $p(z_{t-1}|z_t)$
- Model training can also be done with ELBO in the same way as VAE!



$$\begin{aligned} \log p(\mathbf{x}) &= \log \int p(\mathbf{x}, \mathbf{z}_{1:T}) d\mathbf{z}_{1:T} \\ &= \log \int \frac{p(\mathbf{x}, \mathbf{z}_{1:T}) q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})}{q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})} d\mathbf{z}_{1:T} \\ &= \log \mathbb{E}_{q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})} \left[ \frac{p(\mathbf{x}, \mathbf{z}_{1:T})}{q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})} \right] \\ &\geq \mathbb{E}_{q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})} \left[ \log \frac{p(\mathbf{x}, \mathbf{z}_{1:T})}{q_{\phi}(\mathbf{z}_{1:T} | \mathbf{x})} \right] \end{aligned}$$

**ELBO**

# Diffusion Model

- If  $T = 1$ , it is equivalent to ELBO in vanilla VAE
- The third term is very dominant in training cost!
- Because it has to be calculated  $T$  times and every KL divergence must be optimized simultaneously
- How to be solve?

$$\begin{aligned}
 \log p(\mathbf{x}) &= \log \int p(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} \\
 &= \log \int \frac{p(\mathbf{x}_{0:T}) q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \\
 &= \log \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \\
 &\geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \\
 &= \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{\mathbb{E}_{q(\mathbf{x}_{T-1}|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_{T-1}) \parallel p(\mathbf{x}_T))]}_{\text{prior matching term}} - \underbrace{\sum_{t=1}^{T-1} \mathbb{E}_{q(\mathbf{x}_{t-1}, \mathbf{x}_{t+1}|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_t|\mathbf{x}_{t-1}) \parallel p_\theta(\mathbf{x}_t|\mathbf{x}_{t+1}))]}_{\text{consistency term}} \\
 &= \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \underbrace{\left[ \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))] \right]}_{\text{denoising matching term}} \star
 \end{aligned}$$

# Diffusion Model

- Diffusion model solves this by imposing strong constraints
  1. The latent dimension = the data dimension
  2. Every latent encoder is a linear Gaussian model
  3. The latent in final step  $T$  must be a standard Gaussian
- This cleans up ELBO because encoders  $q(z_t|z_{t-1})$  are parameterless
- An important intuition is that the sum of a Gaussian RV is another Gaussian RV
- i.e.,  $X \sim N(\mu_1, \sigma_1^2), Y \sim N(\mu_2, \sigma_2^2) \Rightarrow X + Y \sim N(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$

# Diffusion Model

- Diffusion model solves this by imposing strong constraints
  1. The latent dimension = the data dimension
  2. Every latent encoder is a linear Gaussian model
  3. The latent in final step  $T$  must be a standard Gaussian
- This cleans up ELBO because encoders  $q(z_t|z_{t-1})$  are parameterless
- Let  $q(z_t|z_{t-1}) = N(z_t; \sqrt{\alpha_t}, (1 - \alpha_t)I)$  for every  $t$
- So  $x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon$  where  $\epsilon \sim N(\epsilon; 0, I)$

# Diffusion Model

$$\begin{aligned}\mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\&= \sqrt{\alpha_t} \left( \sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2}^* \right) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\alpha_t - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2}^* + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\sqrt{\alpha_t - \alpha_t \alpha_{t-1}}^2 + \sqrt{1 - \alpha_t}^2} \boldsymbol{\epsilon}_{t-2} \\&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\alpha_t - \alpha_t \alpha_{t-1} + 1 - \alpha_t} \boldsymbol{\epsilon}_{t-2} \\&= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2} \\&= \dots \\&= \sqrt{\prod_{i=1}^t \alpha_i} \mathbf{x}_0 + \sqrt{1 - \prod_{i=1}^t \alpha_i} \boldsymbol{\epsilon}_0 \\&= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}_0 \\&\sim \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})\end{aligned}$$

- An important intuition is that the sum of a Gaussian RV is another Gaussian RV
- i.e.,  $X \sim N(\mu_1, \sigma_1^2), Y \sim N(\mu_2, \sigma_2^2) \Rightarrow X + Y \sim N(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$



# Diffusion Model

$$\underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \left[ \sum_{t=2}^T \underbrace{\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{denoising matching term}} \right] \star$$

$$\begin{aligned} q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) &= \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)} \\ &= \frac{\mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t}\mathbf{x}_{t-1}, (1 - \alpha_t)\mathbf{I})\mathcal{N}(\mathbf{x}_{t-1}; \sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0, (1 - \bar{\alpha}_{t-1})\mathbf{I})}{\mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})} \\ &\propto \mathcal{N}(\mathbf{x}_{t-1}; \underbrace{\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})\mathbf{x}_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)\mathbf{x}_0}{1 - \bar{\alpha}_t}}_{\mu_q(\mathbf{x}_t, \mathbf{x}_0)}, \underbrace{\frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}\mathbf{I}}_{\Sigma_q(t)}) \end{aligned}$$

# Diffusion Model

$$\underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \left[ \sum_{t=2}^T \underbrace{\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{denoising matching term}} \right] \star$$

?


$$\begin{aligned} \mu_q(\mathbf{x}_t, \mathbf{x}_0) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})\mathbf{x}_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)\mathbf{x}_0}{1 - \bar{\alpha}_t} \\ \Sigma_q(t) &= \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{I} \\ \mu_{\theta}(\mathbf{x}_t, t) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})\mathbf{x}_t + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)\hat{\mathbf{x}}_{\theta}(\mathbf{x}_t, t)}{1 - \bar{\alpha}_t} \end{aligned}$$

$$\begin{aligned} &\arg \min_{\theta} D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)) \\ &= \arg \min_{\theta} D_{\text{KL}}(\mathcal{N}(\mathbf{x}_{t-1}; \mu_q, \Sigma_q(t)) \parallel \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}, \Sigma_q(t))) \\ &= \arg \min_{\theta} \frac{1}{2\sigma_q^2(t)} \frac{\bar{\alpha}_{t-1}(1 - \alpha_t)^2}{(1 - \bar{\alpha}_t)^2} \left[ \|\hat{\mathbf{x}}_{\theta}(\mathbf{x}_t, t) - \mathbf{x}_0\|_2^2 \right] \end{aligned}$$

- Set  $p(\mathbf{x}_{t-1}|\mathbf{x}_t)$  to be Gaussian to reduce KL divergence
- Since the variance is only a function of  $t$ , we just take it

# Diffusion Model

$$\begin{aligned} & \frac{1}{2\sigma_q^2(t)} \frac{\bar{\alpha}_{t-1}(1-\alpha_t)^2}{(1-\bar{\alpha}_t)^2} \left[ \|\hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{x}_t, t) - \mathbf{x}_0\|_2^2 \right] \\ &= \frac{1}{2} \left( \frac{\bar{\alpha}_{t-1}}{1-\bar{\alpha}_{t-1}} - \frac{\bar{\alpha}_t}{1-\bar{\alpha}_t} \right) \left[ \|\hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{x}_t, t) - \mathbf{x}_0\|_2^2 \right] \\ &= \frac{1}{2} (\text{SNR}(t-1) - \text{SNR}(t)) \left[ \|\hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{x}_t, t) - \mathbf{x}_0\|_2^2 \right] \end{aligned}$$

$$\text{SNR}(t) = \frac{\bar{\alpha}_t}{1-\bar{\alpha}_t} = \frac{\mu^2}{\sigma^2}$$


- Finding the appropriate hyperparameter  $\alpha$  can be done with NN
- But this is a bit questionable

$$\text{SNR}(t) = \exp(-\omega_{\boldsymbol{\eta}}(t))$$

$$\frac{\bar{\alpha}_t}{1-\bar{\alpha}_t} = \exp(-\omega_{\boldsymbol{\eta}}(t))$$

$$\therefore \bar{\alpha}_t = \text{sigmoid}(-\omega_{\boldsymbol{\eta}}(t))$$

$$\therefore 1 - \bar{\alpha}_t = \text{sigmoid}(\omega_{\boldsymbol{\eta}}(t))$$

# Diffusion Model

- Expression manipulation never ends...
- We can view the optimization problem of diffusion model in 3 aspects:

- Training original data!

$$\arg \min_{\theta} \frac{1}{2\sigma_q^2(t)} \frac{\bar{\alpha}_{t-1}(1-\alpha_t)^2}{(1-\bar{\alpha}_t)^2} [\|\hat{\mathbf{x}}_{\theta}(\mathbf{x}_t, t) - \mathbf{x}_0\|_2^2]$$

- Training the noise!

$$\arg \min_{\theta} \frac{1}{2\sigma_q^2(t)} \frac{(1-\alpha_t)^2}{(1-\bar{\alpha}_t)\alpha_t} [\|\epsilon_0 - \hat{\epsilon}_{\theta}(\mathbf{x}_t, t)\|_2^2]$$

- Training the score!

$$\arg \min_{\theta} \frac{1}{2\sigma_q^2(t)} \frac{(1-\alpha_t)^2}{\alpha_t} [\|s_{\theta}(\mathbf{x}_t, t) - \nabla \log p(\mathbf{x}_t)\|_2^2]$$

$$\therefore \mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1-\bar{\alpha}_t}\epsilon_0}{\sqrt{\bar{\alpha}_t}}$$

This is the ‘score’!

$$\therefore 1. \quad \mathbb{E}[\mu_z | z] = z + \Sigma_z \left[ \nabla_z \log p(z) \right] \quad \text{“commonly used to correct sample bias”} \quad \text{Tweedie’s Formula}$$

$$2. \quad \mathbf{x}_0 = \frac{\mathbf{x}_t + (1-\bar{\alpha}_t)\nabla \log p(\mathbf{x}_t)}{\sqrt{\bar{\alpha}_t}} = \frac{\mathbf{x}_t - \sqrt{1-\bar{\alpha}_t}\epsilon_0}{\sqrt{\bar{\alpha}_t}}$$

$$\Rightarrow \nabla \log p(\mathbf{x}_t) = -\frac{1}{\sqrt{1-\bar{\alpha}_t}}\epsilon_0$$

# Score-based Generative Model

- The score is the gradient-log of the true data distribution  $\nabla_x \log p(x)$
- That is, for arbitrarily flexible and parameterizable function  $f$  called the energy function, arbitrarily flexible probability distribution:

$$p_{\theta}(x) = \frac{1}{Z_{\theta}} e^{-f_{\theta}(x)}$$

- But normalizing constant  $Z$  can be difficult to compute tractably
- This can be fixed by setting ‘the score’ to gradient-log and finding the score instead

$$\begin{aligned}\nabla_x \log p_{\theta}(x) &= \nabla_x \log\left(\frac{1}{Z_{\theta}} e^{-f_{\theta}(x)}\right) \\ &= \nabla_x \log \frac{1}{Z_{\theta}} + \nabla_x \log e^{-f_{\theta}(x)} \\ &= -\nabla_x f_{\theta}(x) \\ &\approx s_{\theta}(x)\end{aligned}$$

# Score-based Generative Model



- Express the score as a DNN by optimizing the Fisher Divergence

$$\mathbb{E}_{p(\mathbf{x})} \left[ \|s_{\theta}(\mathbf{x}) - \nabla \log p(\mathbf{x})\|_2^2 \right] \quad \dots \text{and this is the vanilla score matching}$$

- The score means the direction the log-likelihood increases
- The Score-based generative model borrows the idea of Langevin dynamics, a molecular system model that can represent molecular diffusion

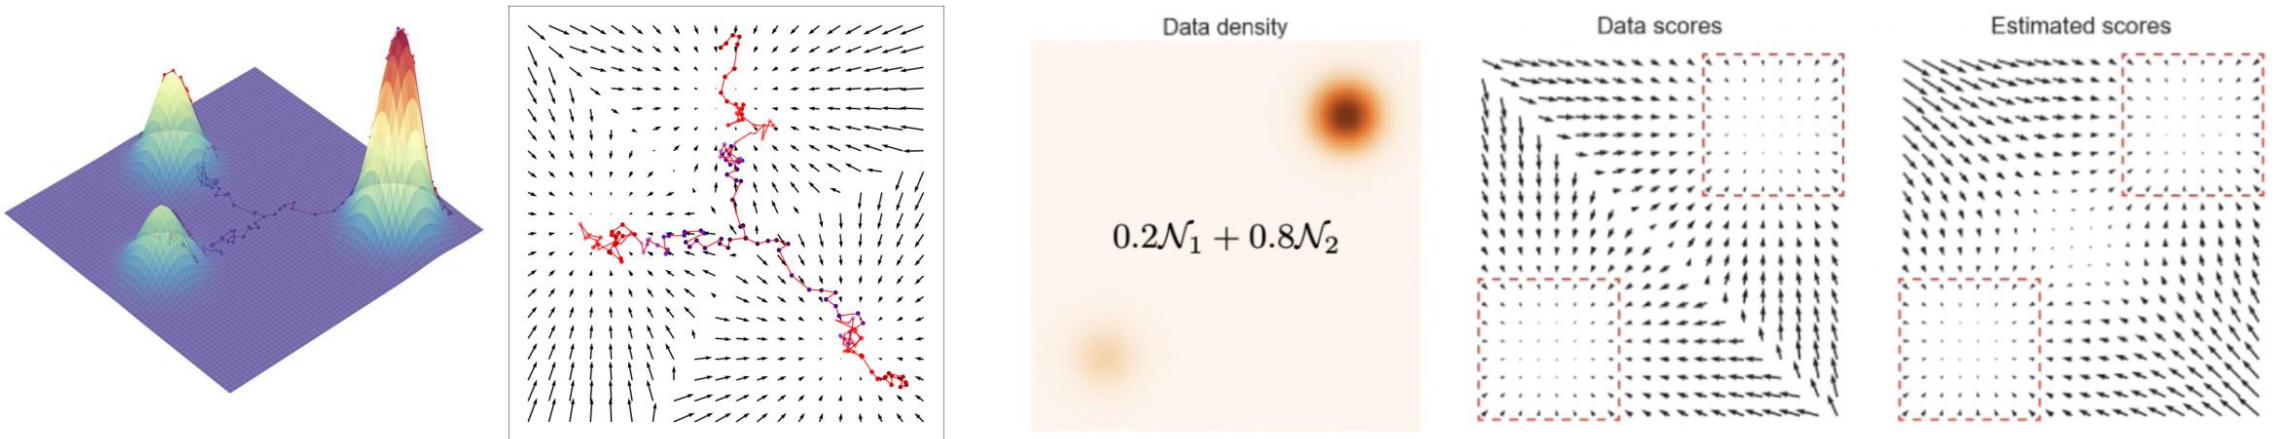
$$M \ddot{\mathbf{X}} = -\nabla U(\mathbf{X}) - \gamma M \dot{\mathbf{X}} + \sqrt{2M\gamma k_B T} \mathbf{R}(t)$$

$$\mathbf{x}_{i+1} \leftarrow \mathbf{x}_i + c \nabla \log p(\mathbf{x}_i) + \sqrt{2c\epsilon} \boldsymbol{\epsilon}, \quad i = 0, 1, \dots, K$$



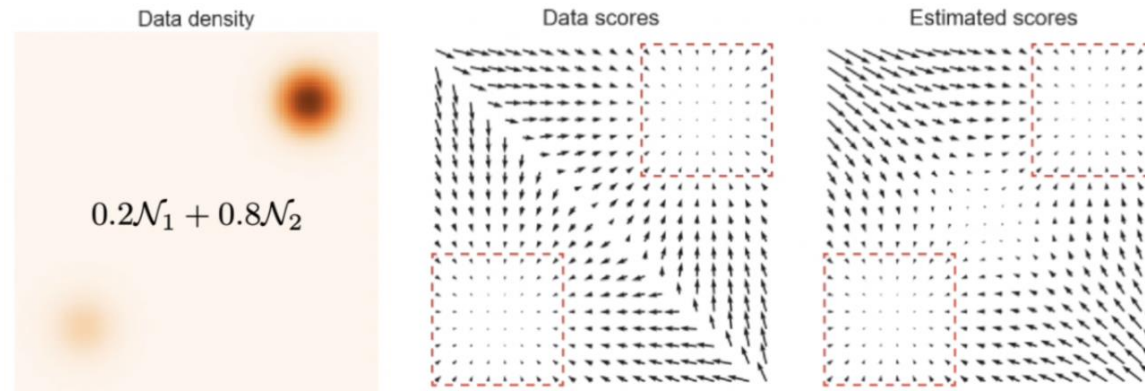
# Score-based Generative Model

- It can estimate an intractable distribution  $p_{\theta}(x) = \frac{1}{Z_{\theta}} e^{-f_{\theta}(x)}$
- But it is very computationally expensive and has various problems:
  1. Ill-defined when  $x$  lies on a low-dimensional manifold
  2. Not be accurate in low density regions
  3. Density is not reflected well (e.g., mixture model) ???



# Score-based Generative Model

- This can be handled by adding Gaussians of different strengths
- The problem of computational cost is solved by estimating the scores of distributions made from noising samples
- Since Gaussian is defined in all spaces, it solves many problems
  1. Ill-defined when  $x$  lies on a low-dimensional manifold **OK**
  2. Not be accurate in low density regions **somewhat ok**
  3. Density is not reflected well (e.g., mixture model) **???**





# Diffusion Model

- Therefore, if we optimize all the Fisher Divergence at each noise level,

$$\arg \min_{\theta} \sum_{t=1}^T \lambda(t) \mathbb{E}_{p_{\sigma_t}(\mathbf{x}_t)} \left[ \|s_{\theta}(\mathbf{x}, t) - \nabla \log p_{\sigma_t}(\mathbf{x}_t)\|_2^2 \right]$$

- A similar expression emerges when solved from the HVAE perspective
- Furthermore, if we modify Langevin dynamics sampling in terms of simulated annealing:
  - Initialize from some fixed prior (e.g., uniform, gaussian)
  - Running Langevin dynamics for each  $t = T, T - 1, T - 2 \dots$
  - The starting point of each step is the ending point of the previous one
- It can be perfectly modeled as a Markovian HVAE!

# Diffusion Model

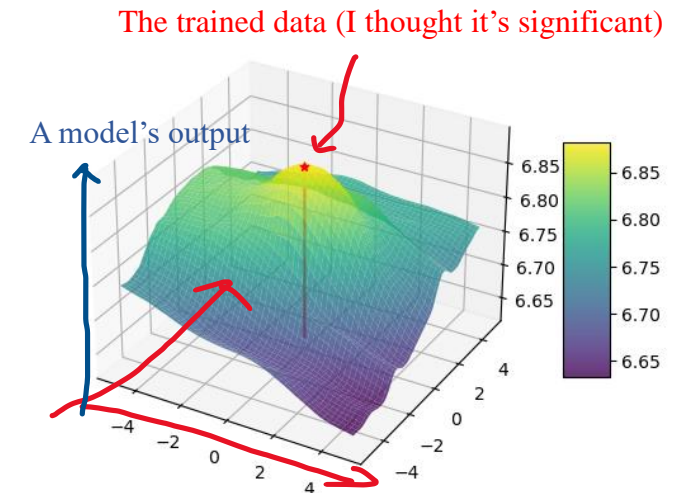
- In conclusion, the diffusion model can be interpreted as an
  1. Hierarchical VAE with strong constraints to deal with the computational cost
  2. Denoising score-based generative model with a clear, comprehensive explanation and great performance
- These are not separate but complementary to each other
- Approaching from different perspectives always gives good ideas!
- I hope this helped you get a rough understanding

# Diffusion Model

- Diffusion models still have very important topics like SDE, but these are not fully understood. (too HARD for me!!!)
- Simple to understand, it seems to be an explanation that unifies the continuous Langevin dynamics process and the discrete DNN process
- The guidance (conditional one) is also possible (e.g., Image-Text)
- However, there are many other stories about these...
- For more details on this presentation, SDE, etc., see below:
  - Luo et al. “Understanding Diffusion Models: A Unified Perspective” [2208.11970]
  - Song et al. “Generative Modeling by Estimating Gradients of the Data Distribution” [1907.05600] (NCSN)
  - Ho et al. “Denoising Diffusion Probabilistic Models” [2006.11239] (DDPM)

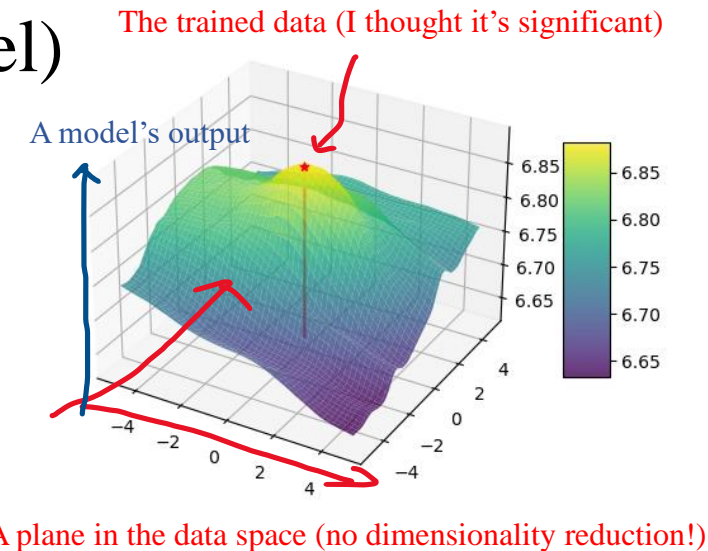
# Experiments and Analysis

- This is one of the results of the experiment in the last presentation
- Simplifying the situation as a classification problem, let us determine whether the model has learned from that data
- Are there any significance features that only trained data have?
- Looking at the picture on the right
- Can we say that the data has been trained?



# Experiments and Analysis


- At first glance, it seems trained (or maybe not)
- Can we set a specific metric?
- My thoughts: Find a feature that is significant compared to other points in the  $L_p$  norm hypersphere centered on the data
- A good candidate is logit (the output of the model)



# Experiments and Analysis


- e.g., if the output of the data itself is noticeably higher than the output of Gaussian noises centered on that data, it has been trained
- Since the model is supervised by training data, the trained data itself will produce a higher output than its vicinity
- But...

If only this cat was trained,  
the output of it would be higher than the other cats  
(this can be seen as overfitting)



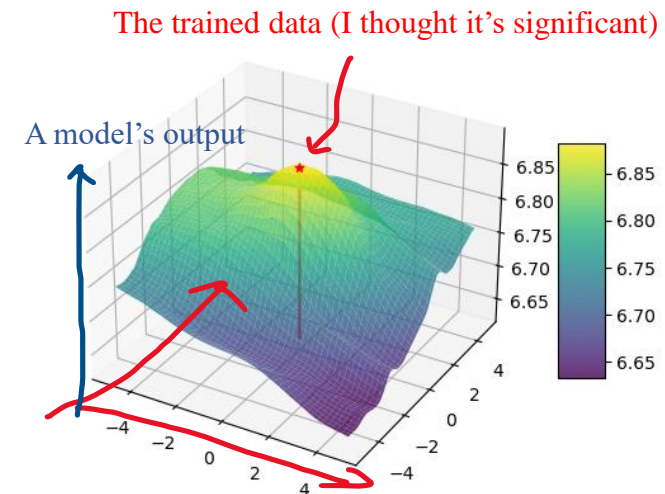

# Experiments and Analysis

- Let us consider a simple binary classification problem
- According to the hypersphere hypothesis, both  $A$  and  $B$  should exhibit significant values (of metric) compared to their surroundings
- But  $B$  will be more significant than  $A$ , since  $B$  is close to the d.b.
- $A$  might even have values same as  $C$ ...



# Experiments and Analysis

- Significant values can be found
- But we do not know if it is caused by just the supervision by that one sample or by others
- No matter how special it is, there is no guarantee that this really happened because of itself



A plane in the data space (no dimensionality reduction!)



# Experiments and Analysis

- However, I believe that the model has prints of the supervision
- In my opinion, the output of the input side of the model with less abstraction is more likely to be affected by the supervision
- But I doubt that this is a good direction
- As we move towards the input side of the model, less meaning is given, and more information is required...

# Experiments and Analysis

- I might change the goal
- Even if I know all parameters of the model, it is difficult to determine
- If the constraint is not a single piece of data, but a small set of data, I can approach it from an adversarial attack perspective
- But it is more meaningful to determine for just one sample

**Thank you for listening**

Presenter: Kim Seung Hwan (overnap@khu.ac.kr)